

ROHITH GOPI

AI Engineer | LLM Integration | RAG Pipelines | Agentic AI | MCP

Chennai | Open to Bangalore | rohiththeg05@gmail.com | +91 73059 19416 | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Final-year B.Tech Computer Engineering (AI/ML) student with production experience building RAG pipelines, MCP servers, LLM-powered applications, and agentic AI workflows, including hands-on work fine-tuning open-source LLMs and building AI agents in a live product environment. Anthropic MCP Certified and AWS Bedrock Certified. Skilled in Python, FastAPI, LLM integration, vector databases, prompt engineering, and on-device inference. Open to relocate to Bangalore for full-time AI Engineer or Machine Learning Engineer roles.

TECHNICAL SKILLS

AI and ML: Large Language Models, RAG Pipelines, Agentic AI, MCP (Model Context Protocol), AI Agents, Agentic Workflows, Prompt Engineering, LLM Integration, Fine-Tuning Open-Source Models, Ollama, ChromaDB, SentenceTransformers, Vector Databases, Deep Learning, NLP, Computer Vision, MLflow

Languages and Frameworks: Python, JavaScript, SQL, Java, FastAPI, TensorFlow, Spring Boot, REST APIs

Cloud and Tools: Azure AI, Azure Machine Learning, AWS Bedrock, Git, GitHub, MLflow, MySQL, Pandas, Jupyter

PROFESSIONAL EXPERIENCE

AI Intern | Zealcre

Jul 2026 - Present | Chennai

- Work as an AI Intern on a software product team, fine-tuning open-source language models to improve task-specific performance and output quality.
- Design and build AI agents that extend the product's automation capabilities, taking features from concept through integration.

AI Engineer Intern | Spotnxt

Mar 2026 - Jun 2026 | Chennai

- Architected and shipped production RAG pipelines and LLM automation features integrated into the core product, owning full delivery from design to deployment in a fast-paced startup environment.
- Collaborated cross-functionally with product and engineering teams to iterate agentic AI features from prototype to production, eliminating manual processes through intelligent automation.

Research Intern - 2D-to-3D Architectural Modeling | SRIHER GitHub

Aug 2024 - Oct 2024 | Chennai

- Collaborated on supervised machine learning pipelines for architectural component detection across 5+ building classes, automating 2D floor plan conversion to structured 3D models.

Research Intern - Trucking Management System | SRIHER GitHub

Mar 2024 - Apr 2024 | Chennai

- Collaborated on an MDP-based truck scheduling engine in Python to optimise fleet utilisation and minimise idle time across dynamic delivery routes.
- Built a multi-role Android application (Drivers, Owners, Admins) with real-time GPS tracking via Java and a Spring Boot backend, eliminating manual dispatch coordination.

PROJECTS

GitHub Intelligence MCP Server [GitHub](#) | [Demo](#)

Python | MCP | Model Context Protocol | Ollama | Qwen 2.5 | [GitHub API](#)

- Architected a production MCP server exposing GitHub-focused AI agent tools including repository discovery, AI-powered code analysis, dynamic tool execution, and on-device LLM explanations via Ollama and Qwen 2.5.
- One of very few freshers in India with hands-on MCP and agentic AI production experience; completed Anthropic's official MCP certification as a direct outcome of this project.

Import-Export Compliance Assistant [GitHub](#) | [Demo](#)

Python | FastAPI | RAG | ChromaDB | SentenceTransformers | BM25 | Cross-Encoder Reranking

- Built a full RAG pipeline combining vector search, BM25 retrieval, and cross-encoder reranking over official policy PDFs, delivering source-attributed responses with zero hallucination through strict prompt and context validation.
- Built a FastAPI backend with a custom chat UI, automating compliance lookups still performed manually across most import and export businesses.

Context-Aware AI Snipping and Summarisation Tool [GitHub](#) | [Demo](#)

JavaScript | Python | Chrome Extension | Ollama | On-Device LLM Inference

- Engineered a Chrome extension with a custom rectangle-based snipping engine that renders AI-generated summaries inline on any webpage including blogs, documentation, and code pages.
- Integrated Ollama for fully private on-device LLM inference with zero data leaving the browser, plus multilingual input, audio playback, and copy functionality.

EDUCATION

B.Tech - Computer Engineering with AI/ML Specialization | SRIHER, Chennai — Expected Jun 2026 | CGPA: 7.5

CERTIFICATIONS

- Anthropic MCP | Introduction to Model Context Protocol | May 2026 | ID: tjeqcyrt8h92
- Anthropic Agent Skills | Introduction to Agent Skills - Claude Code | Mar 2026 | ID: a7z7qxioe99n
- AWS Bedrock | Amazon Bedrock Getting Started | May 2025
- DeepLearning.AI | Generative AI with Large Language Models - Stanford | May 2026 | ID: 4XA5OMONKJ7Z
- Stanford Online | Supervised Machine Learning: Regression and Classification | May 2026
- Microsoft Azure | Azure AI and Machine Learning - 9 Badges | Dec 2025